
A Theory of Group-Based Clientelism

All around the world, politicians harness the material goods they have access to by virtue of being in office in ways that are designed to help them win elections. In some cases, politicians distribute goods in the hope they prove attractive enough to their would-be beneficiaries that they will vote for them on election day. A governing party might enact a programmatic policy that bestows benefits on large classes of voters, such as single parents or the elderly. It might steer infrastructure toward communities containing certain types of voters, such as party supporters, co-ethnics, or people living below the poverty line. It might channel favorable regulatory policies or protection from international competition toward nationally organized groups of voters, such as farmers or local business owners. While the goods in these examples are quite different, the rationale underpinning their distribution in these cases is similar: It is being distributed *in the hope* it moves the needle in favor of the party conferring it in the next election, whether by inducing feelings of loyalty among recipients, creating a perception that the party is competent and cares about the welfare of people like the recipients, or another such mechanism.

In other cases, politicians are not content with the one-sidedness of a relationship in which they deliver goods in the hope they turn into votes. A qualitatively different mode of distribution occurs when a politician is able to *tie* the distribution of a valued good to a person's voting behavior. By "ties," I mean the good is conferred on the understanding that the politician will receive a vote in return, and if it should become apparent that she is not, will not, or did not, the good is withdrawn. Unlike in the former case, in which politicians use the distribution of goods to try to make themselves attractive to voters, the latter case embodies contingency: The distribution of the good is made *contingent* on how a person votes, such that if a vote is not forthcoming or did not materialize, the good is withdrawn.

The contingent mode of distribution is what comparative politics scholars call clientelism. As with the former mode of distribution, non-clientelism, the good changing hands in a clientelistic exchange can be virtually anything, so long as the politician has access to it and the recipient values it. What differentiates the nonclientelistic mode of distribution from the clientelistic mode of distribution is that in the former, voters retain the freedom to pocket the good, whether it be a programmatic policy, infrastructure spending, or favorable regulation, and *vote for someone else* on election day. This freedom is taken away in a clientelistic exchange. In a clientelistic exchange, at least as clientelistic exchanges have traditionally been conceptualized, if a voter wants the good on offer, she has to vote for the politician proffering it. Because her access to the good is tied to her voting behavior and will be cut off if she is found to have voted for anyone other than her politician benefactor, clientelistic exchanges are troubling from the perspective of democratic integrity. They prevent voters from exercising free choice in their decision of who to vote for, and they can keep voters wedded to incumbents who will have few incentives to represent their other, programmatic, interests.

In this book, I make two theoretical claims. The first stems from the observation that the literature on clientelism in political science has been preoccupied with the idea that politicians form clientelistic relationships with *individual voters*. This preoccupation has led scholars to be overly concerned with a particular set of questions, which center on how a politician seeking to buy the vote of an individual voter can consummate the deal to her satisfaction, given that the secret ballot prevents her from being able to observe how this person votes. The secret ballot means that there are hardly any scenarios under which a person's vote choice will be observable to the politician. If a politician receives every vote cast in a given locale, then she will know that everyone there voted for her. If she receives no votes, then she will know that no one did. But outside of those two scenarios, as soon as we are in a world in which the politician receives a *modicum* of votes in a given locale, then anyone who received goods from the politician in exchange for her vote will be able to credibly claim that they supplied one of those votes, and the politician will not know difference. The assumption that politicians' inability to observe how people vote means that they must be using every tool at their disposal to try to ascertain how the people they gave resources to voted on election day has shaped answers to questions ranging from which types of voters are embedded in clientelistic exchanges to what role the politician's intermediaries play in implementing it.

I make the case that under a certain configuration of political institutions, it makes more sense for politicians to form clientelistic relationships

with *groups of voters*. To form this relationship, a politician's electoral district must be divisible into groups of voters, at which levels of electoral support are observable and to which resources are targetable. While clientelistic relationships are possible with any group meeting these conditions, the practice of counting votes at the level of a small, geographically defined group within electoral districts, very common across democracies today, means that *geographically defined* groups are particularly susceptible to being pulled into clientelistic exchanges with politicians. However, to consummate such an exchange, it is not enough for the politician to be able to *observe* disaggregated vote totals at distinct geographic areas within her electoral district; she must also be able to *deliver resources* to those units. The key insight of my theory, group-based clientelism (GBC), is that both conditions must be met: The group has to be able to make its level of electoral support for the politician observable *and* it must be an entity to which resources can be delivered. When both conditions are met, I argue that the ease of relying on publicly available election returns to glean levels of support makes forming clientelistic relationships with groups of voters more attractive to reelection-seeking politicians than forming them with individual voters.

In this chapter, I flesh out this claim. I point out that the possibility that politicians form clientelistic relationships with groups was present in foundational work on clientelism but has been overlooked in recent years, eclipsed by the presumption that politicians prefer to form clientelistic relationships with individuals. Today, almost every study of goods-for-votes exchanges – what the literature calls clientelism – conceives of them as being between politicians and individuals. However, if resource allocations to groups are found to be tied to their voting behavior, meaning that they increase when the group performs “well” for the politician, however defined, and decrease when it performs “badly,” then this is observationally equivalent to an exchange in which an individual's access to resources is tied to his or her voting behavior. As such, this behavior should not be considered *outside* the purview of clientelism. I offer reasons why a politician worried about winning the next election might prefer to form clientelistic relationships with groups over individuals, and I elucidate the conditions under which they will be able to do so. I conclude by taking four longstanding questions of interest in the clientelism literature and explaining how the possibility that politicians are forming clientelistic relationships with groups has the potential to open up new lines of inquiry, which promise to significantly advance our understanding of how politicians use resource allocations to further their electoral interests.

Chapter 4 is devoted to my second theoretical claim, which concerns *how*, exactly, politicians tie group-level allocations to electoral support.

Politicians can do this in several ways. They could set a target vote share for each group and promise to reward any group that meets the target. They could promise to reward any group whose support exceeds that which it delivered to the politician in the previous election. Alternatively, the politician could promise to divvy up any material goods she has access to after the election in proportion to the number of votes she receives from each group, such that a group contributing 20% of the politician's votes would receive 20% of resources, and so on. Drawing on prior work, I explain that when politicians are members of a party that everyone knows is going to win the next election – in other words, a “dominant party” – they will be able to get the most votes, conditional on resources delivered, by converting voting into a competition over which group is the most electorally supportive (Smith and Bueno de Mesquita 2012, 2019; Smith, Bueno de Mesquita and LaGatta 2017; Catalinac, Bueno de Mesquita and Smith 2020).

In such a competition, what this work calls a “tournament,” the incumbent creates the perception that groups will be ranked on the basis of their loyalty in the most recent election, and prizes (in the form of resource allocations) will be awarded on the basis of rank. By offering a very large prize to the highest-ranked group, the incumbent can encourage supportive groups to compete for this position. Among supportive groups, the promise of a very large prize for placing first drives up the incumbent's support, such that rankings can hinge on handfuls of votes. When a handful of votes is enough to change a group's ranking, and a significant increase or decrease in resources awaits any change in rank, voters in these groups have incentives to set aside whatever personal feelings they harbor toward the incumbent or her policies, and think primarily about the influence their vote wields over their group's prize. When resources are structured this way, plenty of voters may dislike the incumbent and her policies but have few incentives to translate that dissatisfaction into their voting behavior. They may *want* another politician elected, but with other people's votes. For the incumbent, structuring resources in this manner increases her support. Moreover, because she uses the *same prize* to motivate support in more than one group, she can deliver less than if she had pursued an alternative scheme through which goods are tied to support.

3.1 WHAT IS CLIENTELISM?

Clientelism has been the subject of sustained attention by comparative politics scholars since as early as the 1970s (Piattoni et al. 2001; Chubb 1982; Shefter 1994, 1977; Scott 1972). Yet it was seminal studies in the

3.1 *What Is Clientelism?*

2000s that catapulted it to the forefront of scholarly inquiry. Kitschelt and Wilkinson (2007) observed that party competition looked very different in countries that democratized later, in the Third Wave, versus countries that democratized earlier, despite the use of similar electoral systems and other political institutions in both. In the late democratizers, parties lacking coherent ideologies, programmatic policies, or records of competence were nevertheless regularly able to attract solid levels of electoral support. In the early democratizers, however, parties of this nature tended not to even exist. The authors attributed this to variation in “citizen–politician linkages,” a concept examined in Kitschelt (2000). Citizen–politician linkages were qualitatively different, the authors argued, in different countries and tended to be remarkably durable over time.

In countries where citizen–politician linkages were what they called programmatic, the authors observed that parties fought elections by coming up with policies targeted at broad swaths of voters in the hope that those policies would prove attractive enough to entice at least some of their intended beneficiaries to vote for them on election day. In contrast, in countries where citizen–politician linkages were deemed “clientelistic,” parties contested elections by using what the authors called direct material inducements – direct payments to voters, or continuing access to employment, goods, or services – to target individuals and small groups of citizens whom the politician felt would be “responsive to such side-payments and willing to surrender their vote for the right price” (Kitschelt and Wilkinson 2007, 2). These inducements were offered on the condition that the voter or small group of citizens vote for the party or politician. Whereas the authors pointed out that programmatic linkages were “indirect” (not targeted at any specific individual voter or small group) and “noncontingent” (there might be a hope that the beneficiary votes for the politician but she does not have to do so), clientelistic linkages were the opposite: “direct” and “contingent.”

Other scholars writing around the same time made similar observations. Stokes (2007) observed that all politicians offered material goods to voters in the expectation that those goods would influence whether or not the intended recipient voted for them. Further, she observed that all types of material goods entailed the redistribution of resources away from larger classes of voters toward smaller classes of voters. A good distributed to a large class of voters on the basis of an ascriptive characteristic such as employment or single parent status was defined as programmatic. A good distributed to a narrower class of voters on the basis of, for example, the type of electoral district in which they resided was defined as pork-barrel politics. Stokes (2007) was primarily

concerned with differentiating these two categories from a third category, which she labeled clientelistic. In a clientelistic mode of redistribution, she pointed out that the sole criterion a politician uses to determine whether to deliver a good to a voter, or whether to end a voter's access to a good she was already enjoying, is "did you (will you) support me?" Whereas in the programmatic and pork-barrel modes of redistribution just described, the criteria was whether or not a voter qualified as a member of said group, in the clientelistic mode of distribution, it was how they voted. A central insight of her study, which has been adopted in virtually every subsequent study of clientelism, is that it is not the *nature* of a good that determines whether or not a politician is distributing it clientelistically, but the *criteria* (the rules) governing its distribution.

Synthesizing some of this work, Hicken (2011) defined clientelism, at least for the purpose of mobilizing votes, as a transaction between politicians (the patrons in the exchange) and clients (voters), in which politicians offer particularistic benefits in exchange for votes. The nature of these benefits were "limited only by politicians' and voters' imaginations, and could range from cash to cookware to corrugated metal" (Hicken 2011, 291). To qualify as clientelistic, he suggested that a transaction had to exhibit two characteristics: "contingency" and "iteration." With regard to the former, he explained that both the politician and voter have to know that the exchange is conditional upon holding up their end of the bargain. The voter has to know that she will only receive the good if she votes for the politician, while the politician has to know that she will only receive the vote if she delivers the good. This "quid pro quo" feature of the clientelistic exchange is what distinguishes it from other forms of particularistic targeting (Hicken 2011, 291).

The second characteristic, iteration, was deemed necessary because exchanges of goods for votes were not, generally speaking, simultaneous. This means that there tends to be a lag between when the voter delivers the vote and when the politician delivers the good (Hicken 2011). This presents opportunities for both sides to defect from the exchange. For the voter, defection means pocketing the good and not voting for the politician on election day, while for the politician, it means receiving the vote but not delivering the good. The promise of an *iterated* exchange is seen to help both sides credibly commit to following through on their promises to the other. More specifically, iteration helps the voter see that, if she receives a good but fails to vote for the politician, the politician will be around long enough to punish her by withholding the good next time. It also helps the politician see that, if she receives a vote from the voter but does not deliver the good, the voter will be around long enough to withhold her vote next time. If the exchange was a single-shot transaction,

3.1 What Is Clientelism?

in contrast, both sides would have powerful incentives to defect and the exchange would be difficult to consummate (Stokes 2005).

What implications does clientelism have for democracy? Stokes et al. (2013) consider this question in depth. They point out that in democracies, almost all individuals over a certain age are given the right to participate in the selection of leaders who will go on to formulate policies that apply to them. The right to the exercise of free choice in this decision of who to vote for is treated as sacrosanct. Elections are critical moments in which the preferences of citizens are translated upwards: Citizens have the right to evaluate, to the best of their knowledge, what their leaders have been doing while in office. If they approve, they can give those leaders a green light to continue; if they disapprove, they have the right to throw their support behind a different set of individuals. The expectation that citizens will be able to use elections to reward or punish their elected representatives for their performance while in office is known as democratic accountability (see also the chapters in Brun and Diamond 2014).

In earlier work, Stokes (2005) made the case that clientelistic exchanges made a “mockery” of democratic accountability. This is because in a clientelistic exchange, it is no longer the politician who is being held accountable by the voter at election time, but the *voter* who is being held accountable by the *politician*. Accountability is thus reversed. Any politician with the ability to punish or reward voters on the basis of their performance at the ballot box will likely be able to subvert (at least partially, if not completely) the ability of voters to use elections in the manner they were intended: as a means of sanctioning politicians for poor behavior while they were in office. In this way, clientelism is thought to undermine voter autonomy. Whenever voters know that they will suffer retaliation for voting for anyone other than their politician benefactor, they will be less likely to do so. This means they are no longer exercising free choice when casting their votes (Stokes et al. 2013).

Extending this logic further, politicians rarely try to buy the votes of *everyone*. The division of the electorate into voters whose votes are being bought and voters whose votes are not is likely to give rise to a fundamental inequality in representation and policymaking, whose consequences for the broader political system are likely to be profound: Namely, the politician benefactor knows that as long as she continues the flow of clientelistic goods to the voters whose votes she is buying (her “clients”), she can afford to ignore their preferences on virtually any other matter (Stokes et al. 2013). Essentially, she is able to bribe her clients to forget their other concerns. Because of the implied threat of retaliation, voters whose votes are being bought cannot credibly threaten to exit the

relationship, which means that they may struggle to convince the politician to pay attention to their other concerns. This creates a situation in which differences in the *type* of voter who is pulled into clientelistic exchanges could be carrying over to quite pronounced biases in public policy. If politicians make a calculated decision to buy votes from the poor, for example, on account of the fact that their votes can be bought with fewer resources, then the risk is that policy will come to reflect the preferences of other, *less* needy groups, which is not socially desirable (see also Diaz-Cayeros, Estévez and Magaloni 2016).

The fact that clientelistic practices are observed in so many countries and pose a threat to democratic integrity helped propel clientelism to the forefront of scholarly inquiry in comparative politics. Central questions of interest have been the types of voters politicians form clientelistic exchanges with and what form those exchanges take (Higashijima and Washida 2024; Yildirim and Kitschelt 2020; Cruz 2019; Bussell 2019; Nichter 2014; Cammett and Issar 2010; Magaloni 2006; Calvo and Murillo 2004; Chandra 2007), whether clientelistic exchanges are always aimed at changing how a person votes (“persuasion”) or whether they are also used to encourage people to turn out to vote (“mobilization”) (Nichter 2018; Rosas, Johnston and Hawkins 2014; Gans-Morse, Mazzuca and Nichter 2014), and what role brokers play in helping politicians consummate their exchanges (Ravanilla, Haim and Hicken 2022; Holland and Palmer-Rubin 2015; Stokes et al. 2013; Larreguy, Marshall and Querubín 2016).

Other studies have tried to identify the conditions under which politicians turn to clientelism to build coalitions of loyal supporters, as opposed to relying on nonconditional modes of distribution (Diaz-Cayeros, Estévez and Magaloni 2016; Weitz-Shapiro 2012; Calvo and Murillo 2004; Wantchekon 2003; Magaloni, Diaz-Cayeros and Estévez 2007), with the aim of shedding light on what brings about the demise of clientelism (Rizzo N.D.; Golden and Nazrullaeva 2023; Kuo 2018; Weitz-Shapiro 2014; Camp, Dixit and Stokes 2014). How clientelism impacts the interests that get represented in politics, and whether clientelism is bad for democracy and economic development and if so, why, have also been the subject of research (Diaz-Cayeros, Estévez and Magaloni 2016; Szwarcberg 2015; Stokes et al. 2013).

3.2 PROBLEM OF CONTINGENCY

Virtually all clientelism research, at least as it pertains to the exchange of goods for votes, whether iterative or single-shot, has following feature. On the one hand, clientelism is deemed worthy of study because of the

3.2 Problem of Contingency

threat it poses to the form and functioning of democracy. As we mentioned, clientelistic exchanges are thought to undermine voter autonomy, limit voter choice, pervert the functioning of elections as a means of holding elected officials accountable, and create unequal representation and a bifurcated electorate (Stokes et al. 2013; Brun and Diamond 2014). Yet on the other hand, clientelism is a practice that, by all accounts, should not be observed in democracies. Why not? In a democracy, the secret ballot affords voters the right to vote in secret (Mares 2015). This means that only in exceptional circumstances – such as when a politician receives every vote or no votes in a locale – is it ever revealed how a given individual voted. The secret ballot gives the people whose votes a politician tries to buy the freedom to pocket the good from the politician, in exchange for a promise to vote for her, but then on election day, renege on this promise and vote for someone else. As long as the politician receives a *modicum* of votes in the voter's locale, then the voter will always be able to credibly claim that *she* was responsible for one of those votes (and is therefore worthy of continued targeting in the next election), and the politician will not have the means of discerning whether or not she is telling the truth.

The fact that politicians cannot verify how people vote means that they have no way of monitoring whether a person with whom they are attempting to cultivate a clientelistic exchange held up their end of the bargain and voted for the politician like they said they would. This creates a fundamental imbalance in the clientelistic exchange: Whereas the *voter* can discern whether or not they receive the good from the politician, the *politician* will not be able to discern whether or not they received the vote from the voter. Without this information, the politician cannot even *identify* clients who reneged on their promises, let alone *punish* the clients who did renege by stopping the flow of resources.¹ And yet, as we explained above, it is the politicians' ability to *retaliate* against clients who renege on their promise, and defect from the goods-for-votes exchange, that is thought to be the glue that binds those clients tightly to them, in ways that undermine voter autonomy, restrict vote choice, and pervert democratic accountability.

The problem of contingency places clientelism scholars in the peculiar position of studying a practice on account of the fact that it poses a threat to democracy, but having to begin by acknowledging that democracies are precisely the kinds of places where clientelistic practices are unlikely to take root. It is not my contention that *no* politician will attempt a clientelistic exchange in a democracy, it is that, by depriving politicians

¹ The exception to this is when a politician receives zero votes in a locale; then she will know that everyone she attempted to buy votes from reneged.

of the means of enforcing such an exchange, the secret ballot should, all else equal, *decrease* the attractiveness of clientelistic practices, relative to alternative ways of soliciting votes. To return to the dichotomy introduced above (between programmatic and clientelistic citizen–politician linkages), we might expect that being unable to enforce the terms of a clientelistic exchange would make electoral strategies on the *programmatic* end of the spectrum more attractive to the politician. By definition, a politician or party pursuing nonclientelistic electoral strategies offer voters policies that they would enact if elected, on the understanding that everyone qualifying for the policy will benefit from it, regardless of whether they voted for the politician.

Given the vast size of the clientelism literature, it is reasonable to surmise that politicians may have found a way around the “contingency problem,” such that they *can* observe individual votes. Have they? The answer appears to be “no.” In a review of all research on clientelism published in comparative politics journals and university press books between 2008 and 2018, which amounts to eighty-two studies, Hicken and Nathan (2020, 280) found that “very few recent studies, if any, document full contingency, with the commitment problem perfectly resolved and no chance for defection by voters.” Of the studies they examined, 41% explicitly rejected the assumption that politicians had to be able to monitor how people voted in order to cultivate an exchange that otherwise met the criteria as clientelistic. Of the remainder, 16% made no claims about monitoring. 38% assumed that the ability to monitor individual votes *was* necessary, but of these, two-thirds presented no evidence of monitoring, while one-third presented evidence that was either from historical cases where the secret ballot was not yet in place or cases in autocratic or semi-autocratic regimes (with one exception, Hungary, but the evidence in that study did not point to widespread monitoring) (Hicken and Nathan 2020, 280–281). Importantly, their survey of the literature did not reveal a single case of the widespread monitoring of vote choices in a modern-day democracy.

The authors argue that the exchanges actually documented by clientelism scholars are better described as “not-quite-clientelism” or “clientelism-adjacent” (Hicken and Nathan 2020, 279). By this, they mean that the exchanges studied hardly ever satisfy the contingency requirement (meaning there were always opportunities for defection), and in some cases, do not even satisfy the iteration requirement (meaning there was little expectation the relationship would continue into the future). Clientelism, the authors report, appears to have become a catch-all phrase applicable to essentially any situation in which politicians are distributing particularistic, excludable benefits to voters around

3.2 Problem of Contingency

the time of elections. As a result, they conclude that “real clientelism – with a clearly contingent, iterated quid pro quo – appears substantially rarer across the world than the size of the literature on clientelism implies” (Hicken and Nathan 2020, 280). Indeed, of the five cases of real clientelism in these eighty-two works, four were in autocratic or semi-autocratic regimes.

One response to this observation might be to double down on the centrality of contingency to the definition of a clientelistic exchange and demand that any exchange in which this condition is not met is given some other label. The authors take a different path. They recommend a reassessment of the centrality of this characteristic to an exchange qualifying as clientelistic. They reason that if a politician without the ability to monitor votes (and thus, enforce the terms of a clientelistic exchange) nevertheless dedicates herself to cultivating one, then the appropriate conclusion to draw is that the politician may not *need* to be able to monitor votes to make clientelism a viable electoral strategy. One reason a politician might choose clientelism even when she cannot verify individual votes is because distributing resources in this fashion buys her credibility in the eyes of voters. If every candidate distributes handouts prior to elections, and evidence suggests that voters like handouts (Wantchekon 2003), then any candidate choosing not to partake in this goody-mongering risks being seen as unconcerned for voters’ welfare. In such a setting, politicians might choose clientelism to avoid the negative signal that they worry opting out of it incurs (Aspinall et al. 2022, 13). They may also choose clientelism, even if they lack the ability to monitor votes, because they see it as insurance against encroachments on their supporters by rival candidates, or because it increases their name recognition in the electorate and builds their brand (Aspinall et al. 2022, 13).

Hicken and Nathan (2020) point out that all electoral strategies involving the tactical allocation of resources with an eye to influencing how people vote can be scaled according to their “yield rate,” meaning the number of votes the politician receives, per dollar spent. At one end of the spectrum is a programmatic electoral strategy. Because politicians promising programmatic policies cannot deny them to voters who qualify for those policies but voted for the opposition, this electoral strategy is expected to have a low yield rate (many people who receive the good will vote for the opposition). At the other end is “real clientelism,” which is expected to have a high yield rate. Because real clientelism is expected to have such a high yield rate, the authors argue that it makes sense for politicians to invest considerable energy and resources into trying to make this strategy work, which amounts to trying to circumvent the secret ballot. To this end, it pays to cultivate intermediaries known as

brokers, whose job is to penetrate social networks in an attempt to identify which types of voters, if bribed, are most likely to follow through on their promise to vote for the politician (Finan and Schechter 2012; Stokes et al. 2013). Recognizing that investing in such surveillance of individuals is costly and decreases the amount of resources available for vote buying, the authors conclude that many politicians are likely to just muddle through, giving particularistic benefits to select voters and making it clear that they expect a vote in return, but anticipating that their yield will be far less than if they had the means of monitoring their votes. By their calculation, however, a “weakly enforced clientelistic exchange” is still likely to yield vastly more votes than the decidedly inferior (from the perspective of yield) programmatic strategy (see also Aspinall et al. 2022).

Golden and Nazrullaeva (2023) also make the point that exchanges need not satisfy the contingency requirement in order to qualify as clientelistic. They define clientelism as “the discretionary use by politicians of resources for electoral purposes” (Golden and Nazrullaeva 2023, 1). Central to this definition is the potential for voters to be excluded from outlays and the explicitly electoral goals of the politician. Critically, however, their definition “omits any notion of exchange, conditionality, or contingency” (Golden and Nazrullaeva 2023, 5). Thus, clientelistic exchanges should be distinguished from other exchanges not by their contingent nature but by their discretionary nature and electoral goals. Thus, clientelism consists of all the ways in which politicians use discretionary resource allocations to win office, whether used defensively (to recoup waning support), proactively (to win over new voters), or as a general signal of their competence, clout, and willingness to work on behalf of certain types of voters. Other studies that adopt definitions that do not embody contingency include Kramon (2017), who defines clientelism as the allocation of material benefits to voters during elections, and Diaz-Cayeros, Estévez and Magaloni (2016, 7), who define it as “discretionary programs targeted to individuals,” which are “unambiguously administered through clientelistic networks.”

3.3 SHIFTING THE TARGET OF THE EXCHANGE FROM INDIVIDUAL TO GROUP

As we have seen, the clientelism literature runs up against the problem that clientelistic exchanges, at least as they have traditionally been conceived, are unenforceable in democracies. Given this, one might think it would make little sense for politicians to make them the core of their electoral strategies. But clientelistic exchanges appear rampant, at least in certain democracies. Faced with this disconnect between what we observe

3.3 *Shifting the Target from Individual to Group*

empirically and the theories that have been developed to account for those empirical observations, recent work posits that we should reconsider the centrality of contingency to our definitions of clientelism. Essentially, this work suggests that we should treat any attempt to strategically confer resources on select sets of voters, with the goal of influencing how those people vote, as “clientelism.” As a result, the field of clientelism scholarship appears poised to move in this direction.

However, there is an alternative path scholars could take, which does not require us to jettison this requirement. Instead, it reinstates contingency as the hallmark of a clientelistic exchange. The reason we may want to treat contingency as special, and demarcate goods-for-votes exchanges in which this condition is met from goods-for-votes exchanges in which it is not met, is because it is the characteristic of a goods-for-votes exchange that poses the gravest threat to democracy. As I explained above, in *noncontingent* exchanges, voters retain the freedom to pocket goods from the politician but vote for someone else on election day. In contingent exchanges, when politicians have the means of monitoring whether the people they gave goods to voted for them on election day, voters lose this freedom and become locked into voting for their politician benefactor. While both types of exchanges confer advantages on politicians with access to resources, and those advantages may also be troubling from the perspective of democratic integrity, a voter’s freedom to cast her vote however she sees fit is more seriously undermined in the latter.

I posit that the reason contingency has loomed so large in the clientelism literature, to the extent that it is labeled a problem, is because scholars tend to view clientelism as a relationship politicians form with *individual voters*. Emblematic of this approach is Weitz-Shapiro (2014, 5–7), who defines clientelism as the “individualized, contingent exchange of goods or services for political support or votes.” She writes that “in a clientelist exchange, benefits are targeted at the individual and linked to individual political behavior” (Weitz-Shapiro 2012, 569). Diaz-Cayeros, Estévez and Magaloni (2016, 8), also, differentiate discretionary programs targeted at individuals, which they call clientelism, from “discretionary social transfers spent on infrastructure, electrification, street pavement, road construction, and so on,” which they call pork-barrel politics. Lindberg and Morrison (2008, 102), too, write: “It is only when there is a promise or implicit agreement about personal favors or goods to be exchanged in return for political loyalty that a clientelistic relationship is established.” More generally, returning to the aforementioned review piece, Hicken and Nathan (2020) found that of the seventy-three articles and nine books on clientelism they surveyed, seventy-eight (95%)

concerned exchanges between politicians and individuals, leaving only four that looked at exchanges between politicians and groups.

It is not hard to see how a focus on individuals leads straight to the realization that, because individual votes are unobservable, clientelism in its purest form is unenforceable. In the literature, this has led to a presumption that politicians must be on a never-ending quest to try to solve this problem and figure out how people vote, and it to this end that much of their energy is devoted. This presumption underpins most of the work done in the strand of clientelism research that focuses on the intermediaries who help politicians implement their clientelistic exchanges (Ravanilla, Haim and Hicken 2022; Bussell 2019; Brierley and Nathan 2021; Mares and Young 2019, 2016; Larreguy, Marshall and Querubín 2016; Finan and Schechter 2012; Holland and Palmer-Rubin 2015; Stokes et al. 2013; Larreguy, Montiel Olea and Querubin 2017; Wang and Kurzman 2007). Known as brokers or vote mobilizers, these intermediaries are presumed to help politicians with two tasks, both designed to help with the contingency problem. One is identifying which voters a politician should forge clientelistic exchanges with. This is thought to require intimate knowledge of the partisan preferences, short-term situational changes (such as a job loss or sick family member), and attitudes toward helping others of the voters in a politician's jurisdiction. The second task is helping the politician discern whether someone who received a good in exchange for a vote followed through and voted for the politician on election day. This, too, is thought to require intimate knowledge about voters. The fact that clientelism saddles politicians with this "problem" (of being unable to observe individual votes) leads to the presumption that intermediaries are employed to help them solve it. Correspondingly, a trove of studies examine how politicians manage these individuals.

However, the foundational literature referenced above adopted a more expansive definition of clientelism. Kitschelt and Wilkinson (2007, 2), for example, theorized that politicians could cultivate clientelistic exchanges with individuals or "small groups of citizens." In the case of the latter, the authors reasoned that the politician could deliver local club goods to groups "who have already delivered or who promise to deliver their electoral support" (Kitschelt and Wilkinson 2007, 10). They argued that group members would be motivated to vote for their politician benefactor in exchange for the local club goods their group would become eligible for. Beyond pointing out that it was decidedly less costly to monitor votes at the level of a group, especially when election results are reported at a fine-grained level such as a polling station, the authors did not theorize as

3.3 *Shifting the Target from Individual to Group*

to the conditions under which politicians might want to form clientelistic relationships with groups over individuals, or vice versa.

Writing in the same volume, Medina and Stokes (2007) argued that when votes are counted at a disaggregated level such as a precinct, politicians could use the promise of large rewards or punishments to induce voters to vote for them. When the reward or punishment was very large, they argued that politicians who cannot observe how individuals vote will nevertheless be able to create the perception that a voter's livelihood will be significantly affected by how they vote. To illustrate this, the authors describe a situation where voters are threatened with paramilitary attacks on their communities if their politician benefactor does not emerge as the top vote-getter in the community. In this scenario, they argue that even though voters will know that the probability that the politician places first will hardly be affected by whether or not they *personally* make the effort to vote, many voters will nevertheless decide that turning out and voting for the politician is worth it to avoid an attack on the community that could kill those close to them.

From the politician's perspective, key to the effectiveness of such a strategy is being able to count votes at a disaggregated level (Medina and Stokes 2007). The authors contrast this with a situation in which returns are counted at a much higher level of aggregation, like a city with a million people. Under these conditions, they explain that such a threat will not be as effective at inducing people to turn out and vote for the politician because each voter will realize that the impact their vote holds over whether or not their politician places first in the city as a whole is negligible. In this setting, the politician will find it harder to create the perception that the welfare and safety of any one voter is meaningfully influenced by how they cast their vote.

Building on these insights, I posit that when vote counts are discernible at the level of groups within a politician's electoral district, and resources are targetable at those same groups, politicians will find it advantageous to kick the clientelistic exchange up a notch, from the level of the individual voter to the level of the group of voters. The advantage in doing so is that the level at which vote tallies are *discernible* is the level at which the terms of any goods-for-votes exchange they cultivate are *enforceable*. Moreover, they are enforceable at virtually no cost to the politician.

To elaborate, consider a politician whose electoral district is divided into a number of geographically defined groups, such as precincts or municipalities, at which votes are counted and tallies reported. After the election, all the politician has to do is consult the public record to glean the number of votes she received from each group. Using these tallies, she will be able to map out the geographic distribution of her support

within the electoral district, with the level of granularity determined by the number of groups into which the electorate is divided and vote tallies are reported. Then, providing she has access to resources that are targetable at the same group, she will have the tools to implement a clientelistic exchange with these groups. By this, I mean that if a group's level of support for her was found to have declined in a given election, she will be able to observe this and retaliate by withholding resources. Importantly, the politician will be able to discern changes in her support from one election to the other without having to divert *any* resources to creating and maintaining a surveillance operation capable of telling her which of the many individuals she attempted to buy votes from voted for her on election day. By setting up the clientelistic exchange with *groups of voters*, how any given *individual voter* casts their ballot is no longer her concern.

In what I am proposing, *group-based clientelism*, the politician's partner in the exchange is no longer the individual but the *group*, and the goal is not to buy the votes of individuals but to buy *support* from groups. This claim differs from existing research on groups in the clientelism literature in subtle yet important ways. One set of studies assume that politicians are trying to buy the votes of individuals, but reason that, because individual votes are unobservable, politicians will come up with ways of using group-level election returns to discern how many individuals they tried to buy votes from followed through and voted for them on election day. Rueda (2017), for example, argues that when a relatively small number of voters are assigned to vote at a given polling station, brokers will set a target number of votes they want for the politician at that station and imply that the flow of goods to the bribed voters assigned to vote therein will be cut off if the target is not met. Because the number of voters assigned to that station is small, bribed voters will be more likely to vote for the politician on account of the fact that their vote wields influence over whether or not their polling station meets its target. Caselli and Falco (2019) make a similar point, arguing that voters behave differently when assigned to small polling stations as opposed to larger ones, due to the greater risk that their vote choice is revealed to politicians. In both of these studies, politicians are thought to be targeting *individuals* (not groups), and mete out rewards and punishments to *individuals* (not groups).

A second set of studies share the assumption that politicians are trying to buy the votes of individuals. To this end, they are presumed to devise ways of using group-level election returns to discern how hard the intermediaries they hire to deliver individual votes are working for them. Larreguy, Marshall and Querubín (2016), for example, argue that local party operatives lack on-the-ground insight into voter preferences and willingness to vote on election day. As a result, they hire brokers capable

3.3 *Shifting the Target from Individual to Group*

of supplying this information. Preventing shirking by their brokers thus emerges as a central concern for the politician. The authors argue that when an electorate is divided into more groups at which vote tallies are reported, politicians will be better able to distinguish between changes in voter turnout that are a function of broker shirking, and changes in voter turnout that are a function of shocks outside of broker control. Leveraging the fact that Mexican electoral precincts gain an additional polling station for every 750 voters, they compare otherwise-similar precincts with and without an extra polling station and find that turnout is higher in precincts with the extra polling station. This is evidence, they argue, that brokers shirk less at smaller polling stations. In this study, the partner in the clientelistic exchange remains the *individual* (not the group), and group-level returns are seen as useful only insofar as they give politicians the tools to monitor how hard others are working to deliver the votes of *individuals* (not groups).

A third set of studies look at the allocation of group-level benefits. It is widely acknowledged that politicians all over the world have access to material goods targetable at groups of voters, whether those groups are geographic-based (such as villages), association-based (such as religious organizations), or sector-based (such as occupational organizations) (Aspinall et al. 2022; Harris and Posner 2019; Hutchcroft et al. 2014; Scheiner 2006; Maclachlan 2006). Politicians in the ruling party are especially blessed when it comes to resources they control. They routinely have access to funds they can use to deliver infrastructure and other projects to geographic areas of their choosing. Because benefits funneled to groups are enjoyed by all group members, they fit the definition of club goods and are variously referred to as “local public goods,” “pork-barrel politics,” or “partisan bias” in the literature (Stokes et al. 2013; Diaz-Cayeros, Estévez and Magaloni 2016).

In contrast to my claim, which is that club goods can be exchanged for electoral support in a clientelistic fashion, a consensus has emerged in the clientelism literature that club goods are analytically distinct from clientelism. Why? Scholars emphasize that while the decision to bestow benefits on certain groups and withhold them from others might be *political*, meaning that it is undertaken with explicitly political (and usually electoral) goals, it is not *contingent*, on the grounds that these benefits cannot be tied to votes in the way that clientelistic goods aimed at individuals are (Diaz-Cayeros, Estévez and Magaloni 2016; Stokes et al. 2013; Weitz-Shapiro 2014; Lindberg and Morrison 2008; Aspinall et al. 2022; Hutchcroft et al. 2014; Hicken 2011).

In what will by now be a familiar refrain, underpinning the claim that group-level allocations lie outside the purview of clientelism is the

assumption that politicians are focused on buying the votes of *individuals*. Because club goods are bestowed on everyone in a group and cannot be withheld from group members who fail to vote for the politician, members receive the (club) good regardless of whether they personally contributed to the quality that made the group eligible for the good. On these grounds, scholars stress that club goods cannot be used to enforce a clientelistic exchange. Lindberg and Morrison (2008, 102), for example, write that club goods such as schools, roads, and electricity for the community are enjoyed by everyone, so cannot be used to enforce patron–client relationships between politicians and individuals. Diaz-Cayeros, Estévez and Magaloni (2016, 8) write that “a key difference between clientelism and the pork barrel is that within a given district politicians cannot screen users with the latter instrument.” Weitz-Shapiro (2014, 7) writes that pork-barreling, defined as the redistribution of funds away from citizens toward residents of select geographic areas, is not tied to how any given individual in the groups votes, and therefore does not fit the definition of “clientelistic” (see also Weitz-Shapiro 2012, 569). In support of this, Aspinall et al. (2022) find that politicians regularly seek to influence allocations of club goods for political ends (this is “meso-particularism” in their language) but rarely make references to any quid pro quo in which the group receiving the good is expected to reciprocate with votes (Hutchcroft et al. 2014). On this basis, the authors conclude that club goods are not being used to buy electoral support; therefore, they are not being used clientelistically. Instead, they argue that the purpose of meso-particularism is to buy the politician credibility, build her brand, and protect her turf from encroachment by rivals.

If we relax the requirement that clientelistic relationships are necessarily between politicians and individuals and devote serious consideration to the possibility that under a certain configuration of political institutions, politicians might be able to form them with *groups*, we can see that the fact that club goods are not targetable at individuals within groups does not negate the possibility they are being used clientelistically. Recently, scholars have begun to investigate the possibility that group-level benefits are used to purchase support from groups, without the presumption that they are used to buy the votes of individuals (Cooperman 2024; Gottlieb and Larreguy 2020; Gottlieb et al. 2019). These studies do not explicitly tout the relationships between group-level support and benefits they observe as “clientelistic.” However, the definition of clientelism offered in the seminal studies discussed above hold that it is not the *nature* of a good that determines whether it is used clientelistically, but the *criteria* governing its distribution. These studies emphasized

3.4 Political Institutions Are Key

that if the sole criteria under which a good is delivered is “did you vote for me?,” then the good is being distributed clientelistically (Hicken 2011; Stokes 2007; Kitschelt and Wilkinson 2007). I argue that if the sole criteria under which a club good is delivered to a group is “how many votes did you return for me?,” then there are few compelling reasons why we cannot also say that the (club) good is being distributed clientelistically.

My claim is not that club goods are always being used clientelistically. Rather, it is that when institutions are configured a certain way, politicians will be able to form clientelistic relationships with groups, in which club goods are exchanged for support from the group as a whole. While *how* the club good is used to buy support from the group will vary according to contextual factors such as the presence of a dominant party, the basic idea is that the politician will make it known that the amount of money to be delivered to each group will be a function of its level of electoral support. It will be made *contingent* on its level of electoral support. If a group’s level of support for a politician drops relative to a previous election or if it is found to be lower than that of other groups (again, we can expect the precise allocation rule to vary according to contextual factors), the politician will penalize it with less money (fewer goods). Conversely, if a group’s level of support for the politician increases relative to a previous election or is found to be higher than that of other groups, the politician can reward it with more goods.²

3.4 POLITICAL INSTITUTIONS ARE KEY

My claim is that politicians can pursue GBC under a certain configuration of political institutions. What is that configuration? GBC is possible when politicians can observe their electoral support at the level of a group within their electoral district and have the ability to target resources at that same level. It is the degree to which the units at which vote shares are observable and resources targetable *overlap*, and are *nested* within broader electoral districts, that determines the viability of this electoral strategy.

Let us unpack this further. When setting up their democratic institutions (or having them set up by someone else), countries have many decisions to make. One of those decisions is where the votes cast in elections will be *counted*. While all countries have voters cast votes at

² While Weitz-Shapiro (2014, 7) argues that politicians’ inability to tie the provision of pork to how an individual within the group votes means that pork is “analytically distinct” from clientelism, she acknowledges in a footnote that “if that degree of contingency is possible, then clientelism, rather than pork is a more accurate description.”

polling stations, countries vary with regard to whether the votes cast in polling stations are counted at the polling station itself or transported somewhere else and counted there. That “somewhere else” could be a central location in the electoral district as a whole, or it could be an entity such as a precinct or municipality located within the electoral district but sitting atop the polling station. When the votes cast at polling stations are transported to a central location at the level of the electoral district and counted there, politicians are usually prevented from being able to observe their vote shares at the level of a group within their district. In this setting, then, GBC will not be a viable electoral strategy. In contrast, when the votes cast at polling stations are counted either at the polling station itself, or are transported to a jurisdiction such as a precinct or municipality located within the electoral district and counted there, the condition that politicians can observe their vote shares at the level of a group within their electoral district is met, and GBC becomes viable.

How many countries enable politicians to observe their vote shares at a disaggregated unit within the broader electoral district? Data on the unit at which votes are counted in 218 countries, both democracies and nondemocracies, was collected by ACE: The Electoral Knowledge Network (2013). Of these, four held no elections, while eight did not make this information available. Of the remaining 206 countries, 178 countries (86% of the total) are coded as counting votes at the polling station or polling center.³ The remaining twenty-eight countries (14%) transport the votes cast in several polling stations to a “special counting center” and count them there.⁴ While a special counting center means that votes cast at polling stations are transported somewhere else to be counted, further investigation was necessary to determine whether the special counting centers used in each country were coterminous with the borders of the electoral district as a whole (ruling out GBC), or were jurisdictions located *within* the electoral district (making GBC viable). Of the twenty-eight countries using special counting centers, the special counting center appears to be coterminous with the borders of an electoral district in four countries: the United Kingdom, Ireland, Namibia, and Botswana.

³ This number includes the 177 countries coded this way in the data, as well as Canada, which counts ordinary votes cast on election day in polling stations (mail-in ballots are counted differently).

⁴ This number includes the twenty-six countries coded as counting votes in special counting centers, as well as the United States and Anguilla, which were coded as “Other” in the data. In the United States, votes are counted in electoral precincts, while in Anguilla, votes are transported to “a place” designated ahead of time for the purpose of counting. Both these cases fit the definition of “special counting center” even if they are not coded as such in this data (this may have been because this terminology is not used in their electoral legislation).

3.4 *Political Institutions Are Key*

In the remaining twenty-four, special counting centers are units within the broader electoral district. Putting together the 177 countries counting votes at polling stations or centers, and the 24 countries counting votes at special counting centers located within broader electoral districts, we find that 202 out of 206 countries (98%) enable politicians to observe their vote shares at a disaggregated unit within the broader electoral district.

In this discussion, I am equating the *counting* of votes at the level of a unit within the broader electoral district with the *reporting* of vote tallies at the same unit. In other words, I am assuming that if votes are counted at the level of a given subdistrict unit, the results of those counts, in terms of votes cast for the politicians and parties competing, will also be available at the same unit. Unfortunately, data on the unit at which vote tallies are reported or made available does not appear to have yet been collected. To gauge the plausibility of this assumption, I narrowed down the 218 countries in the ACE: The Electoral Knowledge Network (2013) data to the 88 democracies and semi-democracies included in the Democratic Accountability and Linkages Project (DALP) (Kitschelt et al. 2013).⁵ Of these eighty-eight countries, eighty-four count votes at a disaggregated level.⁶ A search of the national electoral commission websites in these eighty-eight countries revealed that thirty-five reported disaggregated vote tallies matching the unit at which votes were said to be counted. Of the remaining forty-nine countries, I was unable to ascertain whether the absence of these tallies on commission websites meant that they were not reported, or that they *are* reported, but are simply not published on these websites.⁷

On balance, the assumption that votes counted at a disaggregated level are observable (by politicians at least) at that level is reasonable on two grounds. First, one reason votes are counted at a disaggregated level is to enhance transparency. This suggests that publication of vote tallies is likely to accompany vote counting. Second, politicians are often allowed to send observers to places where votes are counted, also in the name of enhancing transparency and count legitimacy. The fact that the count

⁵ This project collected data from eighty-eight countries in the 2008–2009 period. The criteria for inclusion was that a country had more than 2 million inhabitants and had conducted at least two rounds of national elections under “democratic or semi-democratic conditions.” Five countries not meeting these conditions (Egypt, Indonesia, Malaysia, Pakistan, and Russia) were also included.

⁶ The United Kingdom, Ireland, Namibia, and Botswana were the four that did not.

⁷ In Japan’s case, for example, votes are counted at special vote-counting centers (we explain what those are in Chapter 5) and published by local electoral commissions after elections are held. But with the national electoral commission does not make these counting-center-level vote tallies available on its website.

takes place with witnesses suggests it is not intended to be kept shielded from the eyes of political actors.

Thus, it appears that, in 98% of countries (and 95% of democracies), politicians are able observe their vote shares at a disaggregated unit within their electoral district. In how many of these countries do politicians have access to resources targetable at or around the same level? The vast literature in political science on pork-barreling and the allocation of club goods suggests that controlling government is synonymous with having access to goods targetable at select groups, whether geographically defined or otherwise, at least for countries above a certain threshold of development.⁸ In some countries, being a member of parliament, not necessarily a member of the governing coalition, is a sufficient condition to be able to influence the allocation of such goods. Critically, however, the degree to which the jurisdictions to which such benefits are deliverable *overlap* with the units at which vote shares are observable is a much subtler question, about which we know relatively little.

To elaborate, *perfect overlap* occurs when vote shares are observable at the exact same unit to which central government resources are deliverable and units are perfectly contained within the boundaries of a single electoral district. In Chapter 5, I explain that many politicians seeking election to Japan's all-important Lower House compete in electoral districts with perfect overlap. Votes are counted and reported at the level of the municipality, which is also a jurisdiction to which central government resources are delivered. With some exceptions, municipalities are perfectly contained within a single electoral district. The fact that many municipalities belong exclusively to a *single* electoral district, and politicians can observe their vote tallies in each municipality and influence resource allocations to municipalities, creates a situation where each and every municipality can be pulled into a clientelistic relationship with the politician representing the district. These are *perfect storm conditions* for GBC: Any politician with influence over government resources will be able to tie the amount of resources municipalities receive to the level of support it provided the politician in the most recent election.

In other settings, votes might be counted and reported at the level of a polling station. While central government resources are not normally deliverable to polling stations, polling stations are often perfectly

⁸ For work on pork-barreling, consult Jensenius and Chhibber (2023); Malik (2023); Motolinia (2021); de Kadt and Lieberman (2020); Rickard (2018); Harris and Posner (2019); McMichael (2018); Kramon (2017); Catalinac (2015); Funk and Gathmann (2013); Huber and Ting (2013); Tavits (2009); Golden and Picci (2008); Magaloni (2006); Diaz-Cayeros, Estévez and Magaloni (2016); Hirano (2006); Lee (2004); Horiuchi and Saito (2003); Dahlberg and Johansson (2002); Ward and John (1999) and Dixit and Londregan (1996).

3.4 *Political Institutions Are Key*

contained within the boundaries of a jurisdiction to which central government resources *are* deliverable, such as a town. If those jurisdictions are then perfectly contained within the boundaries of a single electoral district, then we have something akin to perfect overlap, although it requires the politician to take the extra step of *aggregating* their polling station-level vote tallies to a town-level tally. Politicians contesting Lower House elections in Italy and Mexico operate in electoral districts that meet these conditions.

In contrast, *imperfect overlap* occurs when vote shares are either not observable at the same unit to which central government resources are deliverable or they *are* observable at this level, but this unit is not perfectly contained within the boundaries of a single electoral district. When the jurisdictions to which central government resources are deliverable are relatively *large*, for example, they can be coterminous with the boundaries of a single electoral district or encompass several electoral districts. In the former, a politician might be able to observe her vote shares at a disaggregated level within the electoral district, if votes are counted in polling stations, but will not be able to deliver resources to those units (resources are deliverable only to the jurisdiction as a whole, which is coterminous with electoral district). In the latter case, the fact that the jurisdictions to which resources are deliverable contain several electoral districts nested within them mean that politicians do not even have the means of delivering resources to their *own* electoral district, as any resources delivered to the broader jurisdiction containing their electoral district could hypothetically end up in the hands of voters in other electoral districts.

In South Korea, for example, the jurisdiction to which central government resources can be delivered is the municipality. In elections to the National Assembly, three types of electoral districts exist: districts whose borders are coterminous with a single municipality, districts that are part of a larger municipality, and districts containing multiple municipalities. The conditions for GBC are not met in the first two types but are met in the third.⁹ In the case of elections to the House of Commons in the United Kingdom, the central counting system precludes politicians from even viewing their vote shares at a disaggregated unit within their electoral district. Even if they were able to observe disaggregated vote counts, however, the jurisdictions to which central government resources are deliverable (local authorities) are relatively large, with many encompassing more than one electoral district (Fouirnaies and Mutlu-Eren 2015; Ward and John 1999). Thus, the conditions for GBC are not met in elections to the House of Commons.

⁹ Calculated with data collected by myself and Lucia Motolinia.

Other cases of imperfect overlap occur when the units to which central government resources are deliverable might not be particularly large but are nevertheless not contained within the boundaries of a single electoral district and straddle more than one. In Japan, Italy, Mexico, and Ukraine, for example, many politicians operate in settings characterized by perfect overlap (electoral districts are perfectly divisible into units at which vote shares are observable and resources targetable). But some of the politicians in these countries operate in electoral districts in which one or more of the jurisdictions making up the electoral district, at which vote shares are observable and to which resources are targetable, straddle the boundaries of a neighboring electoral district.¹⁰ In these electoral districts, politicians will be able to pull all the perfectly nested jurisdictions into clientelistic relationships with them but would have trouble incorporating the jurisdiction straddling the boundaries of the neighboring electoral district. While the politician would be able to observe her vote shares in the area of that jurisdiction located in her electoral district, she would find it difficult to target resources at that area only. Because resources are targeted at the jurisdiction as a whole, then she would likely find it difficult to ensure that any resources she delivers to the jurisdiction are spent in the section of the jurisdiction located in her electoral district.

A different case of imperfect overlap occurs when the unit to which central government resources are deliverable bears absolutely no resemblance to the unit at which vote shares are observable. Because national governments tend not to administer national elections themselves but *delegate* this task to the lower-tier administrative entities that make up the state, such as municipalities, my research has not unearthed a case like this. Thus, I mention it as a hypothetical. When the lower-tier entities that make up the state are tasked with establishing polling stations, supervising voting, and counting and reporting the votes, then the result is a situation in which vote tallies tend to be observable at the level of this lower-tier administrative entity. A municipality charged with administering national-level elections within its borders would have no reason to establish polling stations or count votes in another municipality. Nor would it have any reason to *mix* the votes cast within its borders with those cast within the borders of other municipalities.¹¹ However, if a national government decided to administered elections directly, without delegating this task to lower-tier units, then theoretically, we might observe a situation in which votes are counted and reported at

¹⁰ Calculated with data collected by myself and Lucia Motolinia.

¹¹ Note that mixing votes is a feature of vote counting in certain countries, such as Japan and Belgium, but this appears to entail the mixing together of votes cast in different polling stations (before they are counted) within the same jurisdiction.

3.4 *Political Institutions Are Key*

units it established, which might not correspond with the borders of the lower-tier units already in existence.

While I was able to establish that politicians can observe their vote shares at a disaggregated level within their electoral district in 95% of democracies, establishing the degree of overlap between those units and the units to which central government resources are deliverable is more difficult, and requires the collection of new data. In gathering it, researchers must be mindful of the possibility that the degree of overlap between the unit at which vote shares are observable and resources deliverable not only varies across country but might also vary within country and over time. As we saw with South Korea, for example, it can be the case that among politicians running for election to the same legislature, some may be running in electoral districts where GBC is possible and others may be running in electoral districts where it is not possible. In Chapter 5, I explain that in the case of Japan, the drawing up of new electoral districts for Lower House elections, combined with municipal amalgamations and urbanization has created variation in the degree to which electoral districts meet these conditions over time. Whereas it used to be the case that virtually every politician competed in electoral districts in which GBC was possible, a growing number of politicians compete in districts where it is not.

Beyond matching up units at which vote shares are observable and resources deliverable, researchers interested in gauging whether GBC is happening in their countries of interest should also consider the following three questions. First, what types of *resources* are available to politicians? The vast literature on pork-barreling suggests that three categories of geographically targetable spending are amenable to serving as fodder for politicians seeking clientelistic exchanges with groups, at least those that are geographically defined.¹² The first category, *project-based funds*, is when the central government takes revenue generated from taxes levied on all citizens and makes a pot of money available for the purpose of funding specific projects in lower-tier administrative units. The units often have to come up with proposals for funding (and sometimes, commit to shouldering part of the cost themselves) and central government bureaucrats weigh up the merits of each project and decide which

¹² In this chapter, I have focused discussion on how politicians could go about cultivating clientelistic relationships with geographically defined groups. It is also possible that politicians could cultivate these relationships with *nongeographically* defined groups, such as sectoral associations. They could do this on the condition that the level of support they receive from each group is observable. Politicians might use goods such as preferential tax rates or regulatory favors to buy the votes of non-geographically defined groups, provided those benefits could be adjusted, upwards or downwards, with an eye to the group's electoral performance.

to fund. Project-based funds are distinct from funds allocated to lower-tier units on the basis of a formula that takes into account need, broadly defined. Several studies tackle the question of whether politicians can influence formula-decided allocations, and while their results are suggestive of influence, their analyses underscore the difficulty of controlling for the myriad factors that go into the formula (McMichael 2017; Ward and John 1999; Banful 2011).¹³

The second category, *supranational funds*, are funds granted to central governments from supranational bodies such as the European Union and the World Bank, on the condition that they are used to create programs designed to alleviate poverty, enhance economic development, improve health outcomes, and so on. It is not implausible that these funds could also serve as fodder for GBC. The European Union Cohesion and Structural Funds, for example, was created to foster convergence in all areas of the EU. Recipient governments receive funds each year, which they are required to use to fund projects at the local level (Dellmuth and Stoffel 2012). While recipient governments must comply with rules in deciding which projects to fund, Muraközy and Telegdy (2016, 324) argue that there is scope for political manipulation, not least because the EU exercises “limited supervision” over this process. Concessionary loans from the World Bank, for example, funded the Ugandan government’s Youth Opportunities Program, which grants funding to help groups of young people in underdeveloped regions start a small business (Blattman, Emeriau, and Fiala 2018).¹⁴

The third category, the *constituency development fund* (CDF), is when a government sets aside a percentage of its annual budget to be allocated directly to members of parliament, for the purpose of funding projects aimed at fostering development in their electoral districts (Malik 2023; Jensenius and Chhibber 2023; Keefer and Khemani 2009; Harris and Posner 2019; Malik 2019). In some of the countries with CDFs, the same

¹³ Project-based funds can include temporary grant programs, such as the local investment grants provided by the Swedish government for investments in ecological sustainable development (Dahlberg and Johansson 2002), and the Australian government’s sports grants (Denemark 2000).

¹⁴ A type of supranational funding that has received sustained attention by political scientists is the conditional cash transfer (CCT). This encompasses cash payments and other in-kind transfers to low-income individuals and families in exchange for ensuring their children attend school, health clinics, and so on (e.g., Conover et al. 2020; Imai, King and Rivera 2019; Correa and Cheibub 2016; Layton and Smith 2015; De La O 2015; Tobias, Sumarto and Moody 2014; Linos 2013; Zucco 2013; Labonne 2012). Because most (although, not all) CCTs consist of money for individuals or families (not groups), these funds may be less likely to serve as fodder for GBC.

3.4 *Political Institutions Are Key*

amount of money is given to all electoral districts, while in others, a formula is used to determine how much each electoral district receives. In still others, electoral districts receive allocations that are in part determined by a need-based formula and in part made up of the same amount all districts receive. The fact that the sitting member of parliament representing the district tends to have complete discretion over how the funds are used within her electoral district, combined with the fact that funds are replenished every year and politicians do not have exert effort to secure them, makes CDFs a near-ideal resource with which politicians could tie the dollar value of projects awarded to groups within the district to their voting behavior.

A second question researchers should consider is: How do politicians exert *influence* over the allocation of these resources? As our discussion of the types of resources that are amenable to GBC has implied, the way influence is exercised will depend on the nature of resources available. In the case of CDFs, politicians gain access by virtue of winning the election. In other settings, tying the amount of resources groups receive to their performance in the most recent election will be less straightforward. While the decision to establish a pot of money that can be used for local development projects and other infrastructure might have been made by politicians, and usually by politicians in the ruling party, decisions about which areas will actually receive the funding are usually delegated to bureaucrats. Rickard (2018, 188) explains why delegating these decisions to bureaucrats helps shield the politician from criticism that she is using these funds for electoral ends and gives her the power to scapegoat bureaucrats for decisions that might later come under criticism. In these settings, we can expect that politicians will seek to influence bureaucratic decision-making via lobbying. In Chapter 5, I explain how in Japan, politicians make calls to bureaucrats on behalf of municipalities and occasionally accompany local politicians to the offices of the relevant government ministry, such that they can make the case for a particular project in person. When the only means of influencing resource allocations is via lobbying, we can expect that politicians will devote more time to lobbying for groups they intend to reward, very little time lobbying for groups they are indifferent to, and the cumulative output of all this lobbying could be allocations that track electoral support.

A third question researchers should consider is: How *dependent* is the lower-tier entity on the funds politicians have to offer? Ultimately, given that the lower-tier jurisdictions to which central government resources are deliverable are local governments, answering this question boils down to unpacking the fiscal relationship between tiers of government in a

given country. Scheiner (2006) made the case that the value lower-tier units will place on the resources controlled by politicians at the national level will be a function of the degree of fiscal autonomy enjoyed by the lower-tier unit. In countries characterized by “fiscal centralization,” the vast majority of lower-tier units could be in a situation where most of their revenue comes from central government transfers, in part because they lack autonomous revenue streams. In countries not characterized by fiscal centralization, lower-tier units can also exhibit high dependence on central government transfers when they lack robust revenue streams, whether due to poor economic conditions, economic shocks, or mismanagement. On balance, higher dependence will mean that a higher value is placed on the government resources proffered by politicians. As a result, the attractiveness of GBC as an electoral strategy, as well as the nature of the groups that are likely to emerge as core supporters of politicians, will likely also be influenced by dependence.¹⁵

3.5 CONTRIBUTIONS

I have argued that under the right configuration of political institutions, politicians will have incentives to pursue GBC, which means they will tie the amount of resources groups in their electoral districts receive to how much support those groups gave them in the most recent election. In Chapter 4, I consider the nuts and bolts of how this can be done. I offer a series of hypotheses centered on how politicians in a dominant party can design allocations to groups such that they lead to the most votes, conditional on resources delivered. The rest of the book turns to the case of Japan to test these hypotheses. In that sense, it returns to the puzzle posed in Chapter 2, about the extraordinary staying power of Japan’s Liberal Democratic Party (LDP), through electoral reforms, transformations to the party system, bumpy economic performance, corruption scandals, and policy mistakes. My results show that GBC has been a central component of the electoral strategies of many LDP politicians since at least 1980, which is when my data begins.

Before doing so, however, let me spell out four ways in which the theory of GBC opens up new avenues of inquiry. These new avenues of

¹⁵ Here, I have discussed the nature of resources available to politicians (the second question) and the degree to which lower-tier units depend on those resources (the third question) as if they are exogenous to the theory. However, in a setting where the institutional conditions for GBC are met, it is possible that politicians find this electoral strategy so attractive that they *create* resources that can be used for this purpose and take actions to prevent municipalities becoming more independent. Under these circumstances, both the nature of resources available and the degree to which lower-tier units depend on them might be partly endogenous.

3.5 Contributions

inquiry have the potential to further our understanding of how politicians use the selective targeting of resources under their control to further their electoral ends, in ways that are likely to have broad consequences for democracy.

3.5.1 *Bringing Institutions Back*

A key finding in the early work on clientelism, which still carries considerable weight in research in the area today, was that formal political institutions mattered relatively little and could not account for why politicians sought to cultivate programmatic versus clientelistic linkage strategies with voters (Kitschelt and Wilkinson 2007). When “other imperatives” made it attractive for politicians to cultivate clientelistic linkage strategies with voters, Kitschelt and Wilkinson (2007, 43) argued that politicians would essentially “work around” formal political institutions such as the electoral system or constitutional structure, such that they would be able to do so. The authors argued that at most, political institutions structured how the exchange was conducted (for example, whether the party was the key actor orchestrating it, or the politician), but not the *demand* for the exchange.

In a systematic treatment of the role of political institutions, Mueller (2007, 259) reached the same conclusion, finding that neither federal, electoral, legislative, nor executive structures appeared to decisively impact the likelihood that a politician would “go clientelistic.” Of the institutions examined, this study found that no specific institution was necessary, nor specific configuration of institutions sufficient, for clientelism to emerge as the dominant linkage strategy in a given country. While institutions might bequeath politicians with incentives to pursue clientelism, he observed that those incentives did not appear “not overwhelmingly strong” (Mueller 2007, 275). Thus, he agreed with the above scholars that the balance of evidence appeared to be that institutions did not influence the demand for clientelism but might influence the form it took. Writing in the same volume, Scheiner (2007, 276) considered the role that Japan’s personalistic electoral system, described in the last chapter, played in pushing Japanese politicians to adopt clientelistic electoral strategies. He concluded that this electoral system was “neither necessary nor sufficient for clientelism in Japan.” Hicken (2007) echoed this claim, arguing that personal vote-oriented electoral systems might increase the probability that politicians would try to buy votes, but did not determine whether or not they did so.

To the extent that the finding that institutions matter less than we think reflects a consensus in the literature, the theory offered in this

book breaks with this consensus. As I explained above, GBC is possible only when institutions are configured a certain way. It is not possible when they are not. This conclusion is important because it suggests that reconfiguring a country's political institutions could eradicate, or at the very least minimize, a politician's ability to pursue GBC. Specifically, taking ballot boxes from all the polling stations within a given electoral district to a central location in the district, and mixing them together first, before votes are counted, would significantly diminish a politician's ability to divide her electoral district up into sections and discern how much support she received from each. To the extent that disaggregated vote counts are preferable on transparency grounds, politicians interested in reform could instead choose to tackle rules governing the drawing up of the boundaries of electoral districts and the relationship between those boundaries and those of lower-tier administrative units. Specifically, doing away with rules that electoral districts must respect the boundaries of existing administrative units, which is another way of saying they must be drawn *around* administrative units, or pursuing amalgamations of those administrative units such that they are larger than a single electoral district or straddle more than one, would also go a long way toward undoing the conditions that make GBC possible.

Early work on clientelism called for scholars to “move beyond the current focus on structures and institutions” (Kitschelt and Wilkinson 2007, 6). Indeed, this call is understandable because the institutions that were the focus of their analyses bore little correlation to the clientelistic practices they were observing. Nevertheless, the seeds of a research agenda focused on the aspect of a country's political institutions identified in my theory, *overlap*, are found throughout the literature on clientelism, in any study that makes use of disaggregated vote counts or documents the prevalence of club goods or transfers targetable at groups (e.g., Cooperman 2024; Gottlieb and Larreguy 2020; Gottlieb et al. 2019; Rueda 2017; Medina and Stokes 2007; Scheiner 2006). The findings of this book lead me to urge anyone interested in the question of how politicians engage in the selective distribution of benefits around the time of elections to consider whether the institutions in their country of interest are conducive to GBC, and if so, to collect evidence that would enable them to evaluate whether or not it is a component of politicians' electoral strategies. Once this research has been carried out on numerous countries, we will have the tools to examine whether, conditional on meeting the conditions for GBC (*overlap*), it is always pursued by politicians, or whether the degree to which it plays a role in politicians' electoral strategies is contingent on the presence of other institutions, such as a particular electoral system.

3.5 Contributions

3.5.2 *New Ideas about the Role of Brokers*

The main conclusion emanating from the vast sub-genre of clientelism research that looks at brokers is that in countries where clientelistic practices are observed, politicians are nowhere without their brokers (e.g., Ravanilla, Haim and Hicken 2022; Brierley and Nathan 2021; Larreguy, Marshall and Querubín 2016; Holland and Palmer-Rubin 2015; Stokes et al. 2013; Baldwin 2013). Brokers are thought to play a critical role helping politicians sustain their clientelistic electoral strategies, first by helping them identify who to buy votes from and second by helping them discern which voters they bought votes from actually held up their end of bargain and voted for the politician on election day.

When politicians are pursuing GBC, however, how a given individual within their electoral district voted will no longer be their primary concern. Because politicians are buying support from groups, not votes from individuals, I expect they will have less of a need for the intimate knowledge and surveillance of individual voters that brokers in many settings are presumed to carry out. When politicians are pursuing GBC, their main concerns will likely boil down to how to create an equilibrium in which voters understand that the amount of resources their group receives will be contingent on its level of support for the politician, on the one hand, and how to harness the resources available to them to generate bite-sized goodies that can be channeled at the groups in their electoral district, on the other. In a democracy, creating the perception that one's access to resources is contingent on how one votes has to be done in subtle ways that avoid raising red flags that the integrity of voting is being undermined. In settings where GBC is occurring, then we expect that the main function of brokers will be to help politicians with these two tasks. It is perhaps the settings in which politicians are *unable* to pursue GBC – again, this is likely to be a function of the configuration of political institutions – that the main function of brokers is to anticipate and monitor individual votes.

In support of this, recent research has unearthed findings that are difficult to square with the claim that brokers are employed mainly for their ability to anticipate and monitor individual votes. Brierley and Nathan (2021) conducted a large-scale voter survey in Ghana. Some of their respondents were employed as brokers, while some were not. Comparing brokers with non-brokers, they found that brokers only had slightly more information about the voting intentions of their fellow voters than non-brokers. Moreover, brokers did not have noticeably deeper ties with their fellow voters than non-brokers and were not more politically active during or after election campaigns than non-brokers. On balance, they

found that if anything, brokers appeared to have stronger ties *upward* (to local elites, defined as local party leaders, bureaucrats, and politicians), rather than *downward* (to voters). They acknowledge that their findings are hard to reconcile with the idea that brokers are agents of *politicians*, helping them anticipate and monitor votes. Instead, they suggest that, at least in the Ghanaian context, brokers might be agents of *voters*, helping them identify and secure resources from their politician benefactor.¹⁶

Baldwin (2013)'s research in Zambia lends further credence to the possibility that brokers might be helping with the task of *delivering goods*, not with the task of inferring individual votes. She argues that in settings characterized by low state capacity, voters attach great value to politicians who can bring home central government resources in the form of infrastructure projects and the like. In these settings, she posits that the main function of brokers ("patrons," in her terminology) is to help politicians bring to fruition the projects for which they secure funding.¹⁷ The help brokers provide likely ranges from identifying problems in the local community that need to be addressed, organizing local resources, and ensuring buy in from the community for projects. In short, brokers "serve as the technology by which local goods and services are demanded and delivered" (Baldwin 2013, 795). She reasons, in turn, that voters are more likely to believe that their elected representatives will be able to bring funds to the community if those representatives have good ties with brokers. She finds that elections tend to turn on which candidates have managed to secure the backing of the broker, and winning candidates with that backing end up providing more local public goods.

The theory offered in this book helps us see that the role of brokers may be systematically different in settings where GBC is possible, relative to settings where it is not possible. Once we have established that politicians are engaging in GBC, future scholarship should subject this hypothesis to rigorous empirical scrutiny. The fact that all three types of the club goods politicians generally have access to require the identification of projects in a community, for which funds can be granted, and two require communities to put together proposals for projects ahead of time (project-based funds and supra-national funds, respectively), imply that elected representatives will need *someone* to help them with this task. That someone may be the broker.

¹⁶ However, Ravanilla, Haim and Hicken (2022) conducted a similar survey in the Philippines and found evidence that was more in line with the possibility that brokers are agents of politicians.

¹⁷ Essentially, patrons play the role that bureaucrats might play in a setting with higher state capacity.

3.5 Contributions

3.5.3 *Clientelism Occurring at High Levels of Development*

The assumption that politicians are buying votes from individuals, combined with the observation that politicians do not try to buy votes from all individuals, led scholars to ponder the question of which types of voters politicians try to pull into clientelistic exchanges. This produced an early consensus that poverty was a critical factor (Brusco, Nazareno, and Stokes 2004; Chubb 1982). From a politician's perspective, the weight of a poor person's vote is equivalent to the weight of a wealthy person's vote, but cheaper to buy. The politician knows that the poor voter is likely to surrender her vote for less (in terms of resources) than the wealthy voter, making it far more expedient to target the former (Hicken 2011).

It follows that as economic development proceeds and incomes rise, voters will require more and more resources in exchange for their vote and may even develop an antipathy toward clientelism on account of its perceived damage to the integrity of their country's democracy (Weitz-Shapiro 2014). Businesses, too, can be expected to start resenting the corrupt practices associated with clientelism and begin to lend their support to politicians championing reforms to clean up the state apparatuses facilitating it (Kuo 2018). At a certain point, the literature suggests that politicians will decide to walk away from clientelism in favor of alternative electoral strategies that do not encompass such explicit *quid pro quos*. In support of this, research shows that within countries, it is the poorer voters who are more likely to end up in clientelistic exchanges with their politicians, and across countries, it is the poorer countries that supply the bulk of evidence of clientelism (e.g., Golden and Nazrul-laeva 2023; Kitschelt and Wilkinson 2007; Calvo and Murillo 2004; Scott 1972).

As I have explained, GBC is an electoral strategy, which means it is designed to help politicians win elections and reduce uncertainty surrounding reelection. It is safe to assume that these goals are shared by politicians everywhere. In my theory, the main factor determining whether a given politician will alight on GBC as a means of mobilizing votes is whether the political institutions under which she is competing afford her the ability to observe her vote share at a disaggregated level within her electoral district and deliver resources to that same level. In principle, then, providing that these institutional conditions are met, GBC is possible at varying levels of economic development, provided that the pool of resources controlled by the central government and disbursed by politicians in exchange for electoral support is sizeable enough to be of interest to the jurisdictions at which votes are counted and resources targetable.

At low levels of economic development, resources commanded by the central government will likely be of interest to almost all of the lower-tier administrative entities. As countries become wealthier, the degree of attractiveness of those resources will become a function of their value, relative to the value of the administrative entity's autonomous revenue. Looking around the world, we can see plenty of examples of countries at the highest echelons of economic development that nevertheless have a relatively large number of lower-tier administrative entities that lack robust revenue streams, whether due to a weak local economy, rules surrounding taxation, or economic shocks such as natural disasters or mismanagement. While the question of how politicians will go about tying the amount of resources groups receive to their voting behavior is the subject of Chapter 4, we expect that when politicians are cultivating clientelistic relationships with groups, groups with the *least* robust autonomous revenue sources will be disproportionately represented among the groups at the *highest* echelons of electoral support for the politician.

Viewed in this light, then, the theory of GBC expects that shifts in a country's wealth, international standing, or technological prowess will not be sufficient to shift politicians away from this electoral strategy, unless they are accompanied either by reforms that end politicians' ability to observe their vote shares at a disaggregated unit within the electoral district and target resources to those units, or a situation where the lower-tier units become so wealthy that the resources commanded by the central government are no longer of interest. This, too, represents a proposition that should be subject to rigorous empirical scrutiny in future work.

3.5.4 Democracy Is Undermined, But by Club Goods

Studies widely credited for laying the foundation for the modern-day study of clientelism did not conceive of it as a relationship limited to politicians and individual voters (e.g., Kitschelt and Wilkinson 2007; Medina and Stokes 2007; Hicken 2011). Nevertheless, somewhere along the way, scholars began to restrict their focus, theoretically and empirically, to individual-based clientelism. In doing so, they ran up against the reality that virtually all democracies today safeguard the right of voters to cast their ballots in secret. This means that politicians cannot enforce the terms of any clientelistic bargain struck with individuals. Because individuals can always renege on such a deal, politicians can shower material goods on individuals during election campaigns, and they can do so with an eye to influencing how those individuals vote, but they will not be able to even *discern* whether a given individual has held up her end of the bargain and voted for the politician, let alone *retaliate* against anyone

3.5 Contributions

who reneges. Thus, the practice of showering material goods on individuals during election campaigns, documented all over the world in a variety of capacities, does not, strictly speaking, satisfy the requirements of clientelism.

One might think there is little to be gained from quibbling over definitions. Why does it matter if a practice that does not fit the definition of clientelism continues to be given this moniker? It matters because the exchange of material goods for electoral support poses a graver threat to the integrity of a country's democracy than the practice of treating voters with material largess in the *hope* this influences how they vote. The latter practice might unfairly advantage politicians with more largess to deliver and could end up amounting to a colossal waste of government resources, but fundamentally, the recipients of that largess retain the right to rebuff the politician's attempts to influence them and will suffer few consequence for doing so. This is not the case with the former practice. With real clientelism, voters know that they will only receive certain resources or avoid having their resources taken away if they cast their vote correctly. This means that its consequences for democracy and democratic accountability are more troubling. For this reason, it is important to delineate a politician's attempt to *tie* the provision of material goods to how people vote (the original definition of clientelism, which embodies contingency) from a politician's attempt to use material goods to *try to influence* how people vote (the more catch-all definition of clientelism in recent work).

Once we reinstate contingency as *the* component that differentiates electoral strategies based on the selective targeting of material goods, we can see that attempts to *buy* levels of electoral support from groups in a politician's district (GBC) poses a graver threat to democracy than attempts to *treat* individuals with material goods in the hope that this wins the politician more votes (the "clientelism-adjacent" strategies characteristic of the clientelism literature today). By extension, this means that a whole category of material goods that scholars have assumed posed *less* of a threat to democracy – club goods – may be posing a much bigger threat than we realized. The flip side of this is that the mode of distribution scholars have been most troubled by, essentially any good given to an individual by a politician who seeks a vote, may be posing less of a threat than we realized. To this end, we urge future scholars to devote more attention to the possibility that, conditional on their country of interest meeting the conditions for GBC, politicians are pursuing it. This will entail investigating whether club goods are being used clientelistically. We also urge future scholars to engage in a more considered comparison of the downsides of all electoral strategies involving the use of material goods to influence votes, those embodying contingency and those not.